

# Q4 - Bayesian Anscombe


Varun Nayyar

30/04/2019

Fit linear regression's to anscombe's quartet to get good fits (type 'anscombe' into R). - Can you choose good prior's and models to get good fits?

Let's plot the anscombe quartet

```
par(mfrow=c(2,2))
p1 = ggplot(anscombe, aes(x1, y1)) + geom_point()
p2 = ggplot(anscombe, aes(x2, y2)) + geom_point()
p3 = ggplot(anscombe, aes(x3, y3)) + geom_point()
p4 = ggplot(anscombe, aes(x4, y4)) + geom_point()
grid.arrange(p1, p2, p3, p4)
```



This is famous for having the exact same linear regression despite having very different true underlying data. I'm going to do an example on quartet2. I recommend trying different versions and seeing the effect on the posteriors

```
# copied from McElreath
library(rethinking)
```

```
## Loading required package: rstan
```

```


## Loading required package: StanHeaders
## rstan (Version 2.18.2, GitRev: 2e1f913d3ca3)
## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)
## Loading required package: parallel
## rethinking (Version 1.80)
mN <- quap(
  alist(
    y2 ~ dnorm( mu , sigma ) ,
    mu <- a + b* normx,
    normx <- x2 - mean(x2),
    a ~ dnorm( 3 , 3 ) ,
    b ~ dnorm( 0.5 , 3 ) ,
    sigma ~ dunif( 0 , 3 )
  ), data=anscombe )

post = extract.samples( mN , n=20 )
precis(mN)



##           mean          sd      5.5%    94.5%
## a      7.4445409 0.3365698 6.9066373 7.982444
## b      0.5000002 0.1067699 0.3293614 0.670639
## sigma 1.1205215 0.2394945 0.7377630 1.503280

plot(anscombe$x2, anscombe$y2,
      ylim=c(2, 11))
#plot(p2)
for (i in 1:20){
  curve( post$a[i] + post$b[i]*(x-mean(anscombe$x1)) ,add=TRUE, col=col.alpha("steelblue",0.3))
}

```



```
post2 = extract.samples( mN , n=1000 )
ap = ggplot(post2) + geom_histogram(aes(a), bins=46, fill="steelblue")
bp = ggplot(post2) + geom_histogram(aes(b), bins=46, fill="firebrick")
grid.arrange(ap, bp, nrow=1)
```



- Show what happens when the quadratic term is added to the 2nd quartet on the posterior of the other terms.

```
ansfilter = anscombe %>% transmute(y2, normx2 = x2-mean(x2), x2sq=x2^2) %>% mutate(normx2sq = (x2sq-mean(x2sq))^2)

mQ <- quap(
  alist(
    y2 ~ dnorm( mu , sigma ) ,
    mu <- a + b * normx2 + c * normx2sq ,

    a ~ dnorm( 3, 3 ) ,
    b ~ dnorm( 0 , 3 ) ,
    c ~ dnorm(0, 1),
    sigma ~ dunif( 0 , 3 )
  ), data=ansfilter,
  #start = list(a=7.5, b=2.78, c=-0.1267),
  #start = list(a=7.5, b=0.5, c=-0.1267),
  debug=FALSE)


precis(mQ)
```

##	mean	sd	5.5%	94.5%
## a	7.4684708	0.2554018	7.0602894	7.87665226
## b	1.0536122	0.3698273	0.4625567	1.64466775
## c	-1.8606623	1.2110735	-3.7961916	0.07486698
## sigma	0.8477207	0.2534593	0.4426438	1.25279761

```

postQ = extract.samples( mQ , n=1000 )
ap = ggplot(postQ) + geom_histogram(aes(a), bins=46, fill="steelblue")
bp = ggplot(postQ) + geom_histogram(aes(b), bins=46, fill="firebrick")
cp = ggplot(postQ) + geom_histogram(aes(c), bins=46, fill="seagreen")
grid.arrange(ap, bp, cp, nrow=1)

```



- Use a T distribution on the error model for the 3rd quartet, one with a low df and another with high df. Left as an exercise for the reader Use a `dt` instead of a `dnorm`. Note the `ncp` parameter and ordering.

- Invert y and x for Anscombe's 4th and model. Can you make the prior strong enough to overcome the outlier?

Left as an exercise for the reader. Play with the various priors on the `slope` and see what gives more robustness. How much information are you adding to this system?

- (Bonus) Without using quap, can you create a naive linear regression using something like rejection sampling?

```

x = anscombe$x1
y = anscombe$y1

Nsample = 10000
a = rnorm(Nsample, 3, 3)
b = rnorm(Nsample, 0.5, 1)

outval = rep(0, Nsample)
# we assume that the datalen is much smaller

```

```

# than the number of parameters to try
for (i in 1:length(x)){
  pred = a + b * x[i]
  # we use the log likelihood since
  # prod(probabilities) = sum(logprobs)
  # and we have less underflow this way
  outval = outval + dnorm(y[i], mean = pred, sd = 1, log = T)
}
# normalise and turn into probabilities
probs = (outval - max(outval)) %>% exp

```

Using a technique called importance sampling, I draw a weighted histogram. Effectively each sample is proportional to it's probability. Hence things closer to the centre are worth more

```


library(plotrix)
weighted.hist(a, probs, seq(0,6,0.4))

```

```

## Warning in weighted.hist(a, probs, seq(0, 6, 0.4)): Not all values will be
## included in the histogram

```



```
sum(probs)
```

```
## [1] 95.80634
```

For a 1000 sample, the above is a measure of how many samples worth of inference I've done